



Language metrics show a much smaller separation

QA/open response: 4,628 examples | MCQ-labeled diagnostic subset: 4,905 examples

QA + open-response token F1



Answer-content matching · diagnostic only



WARNING: MCQ choices were omitted in training and inference; answer matching is not standard MCQ accuracy.

Token F1 measures lexical overlap, not clinical correctness. One seed per scale; no uncertainty intervals.