



Language metrics show a much smaller separation

Ordinary QA/open: 4,461 examples | Converted open-QA: 4,938 examples

QA + open-response token F1



Converted open-QA exact match



Corrected answer-text targets; no candidate options shown. Open-QA exact match is not MCQ accuracy.

Token F1 measures lexical overlap, not clinical correctness. One seed per scale; no uncertainty intervals.